


Exercise: Meta Learning for Static Regression

The “sine” meta learning example

- Sines of varying phase and amplitude

$$y(x) = A \sin(x + \psi)$$

- Simple dependency, but just $K = 5$ points per dataset!
- Taken from the famous MAML paper
- The “Hello World” Meta Learning problem!



Meta Learning Setting

- We have access to a *meta dataset*, namely a collection (possibly unlimited) of datasets

$$\mathcal{D} = \{D_1, D_2, \dots, \}$$

- In the case of static regression, each dataset D_i is an unordered collection of K input-output pairs.

$$D_i = \{(x_{i,1}, y_{i,1}), (x_{i,2}, y_{i,2}), \dots, (x_{i,K}, y_{i,K})\}, \quad x_{i,j} \in \mathbb{R}^{n_x}, y_{i,j} \in \mathbb{R}^{n_y}$$

- The datasets D_i are assumed to be *similar* to each other. They are thought as realizations from a probability distribution $p(D)$
- Each dataset D_i is split in a train and a test portion:

$$\mathcal{D} = \{(D_1^{\text{tr}}, D_1^{\text{te}}), (D_2^{\text{tr}}, D_2^{\text{te}}), \dots\}$$

Meta Learning: In-context learning

- We define an **meta-model** $\hat{y} = \mathcal{M}_\phi(x, D)$ which takes as input:
 - a test point x
 - a training dataset D (i.e., K input/output data points).

It produces as output the prediction $\hat{y}$ corresponding to input x , given as a *context* the dataset D .

- Implicitly, the meta-model estimates a *system-specific* model on D and applies it to x
- We train the meta model by minimizing over the meta-dataset:

$$J(\phi) = \sum_{i=1}^b \sum_{j=1}^K \mathcal{L}(y_{i,j}^{\text{te}}, \mathcal{M}_\phi(D_i^{\text{tr}}, x_{i,j}^{\text{te}})).$$

- The meta model learns to predict the output with just K input/output examples


In-context learning architecture for static regression

- The meta-model $\hat{y} = \mathcal{M}_\phi(D, x)$ has structure:

$$c = \text{DeepSets}_{\phi_1}(D)$$


$$\hat{y} = \text{MLP}_{\phi_2}(c, x).$$

DeepSets is *invariant to permutations* of the K input/output examples in D




Performance with meta learning (in-context learning)

Sampled datasets



The meta model learns to make pretty good predictions with $K = 5$ data points

In-context learning architecture for system identification



- May be seen as a simplification of the encoder-decoder Transformer for system identification

$$\hat{y}_{m+1:N} = \mathcal{M}_{\phi}(u_{m+1:N}, u_{1:m}, y_{1:m})$$

- The Transformer's attention layers are ideal to deal with intricate temporal dependencies of SYSID.
- For the sines example, the simple DeepSet+MLP architecture is more than enough...

Exercise


- Implementation of the in-context learning architecture for static regression DeepSets + MLP is sketched in `in_context_learning_sketch.ipynb`
- Use it to tackle the meta learning exercise (`meta_learning_exercise.ipynb`)
- Optionally, apply hyper-networks and MAML (sketched in `maml_hypernet_sketch.ipynb`, described in the next slides...)

Meta Learning with Hyper-Networks

- A hyper-network provides weight and biases of the model MLP describing the relation $x \rightarrow y$

$$\theta = \text{DeepSets}_{\phi}(D)$$

$$\hat{y} = \text{MLP}_{\theta}(x).$$



- The hyper-network may be seen as a *learned* system identification algorithm

Gradient-Based Meta Learning (MAML)

- The learned algorithm has much more structure: it is one step of gradient descent!

$$\begin{aligned}\theta &= \text{Alg}_\phi(D) = \phi - \alpha \nabla_\phi \mathcal{L}(\phi, D), \\ \hat{y} &= \text{MLP}_\theta(x).\end{aligned}$$

- We learn the best parameter ϕ such that, with one (or more) steps of gradient descent on the training dataset, performance measured on the test datasets is good.

$$J(\phi) = \sum_{i=1}^b \mathcal{L}(\text{Alg}_\phi(D_i^{\text{tr}}), D_i^{\text{te}})$$

Finn, Chelsea, Pieter Abbeel, and Sergey Levine. "Model-agnostic meta-learning for fast adaptation of deep networks." International conference on machine learning. PMLR, 2017.



INTERNATIONAL FEDERATION
OF AUTOMATIC CONTROL



21st IFAC Symposium on System Identification SysID2027

July 7-9 2027, Lyon, France



International Program Committee

IPC Chairs

Maarten Schoukens,
Marco Forgione

Editor

Kévin Colin

National Organizing Committee

NOC Chairs

Xavier Bombois
Marion Gilson,
Guillaume Mercère

Important Dates (Deadlines)

Submission site opens Jun. 01, 2026

Joint journal/conference
submission papers Oct. 01, 2026

Invited session,
Open Invited Track
& Tutorial proposals Oct. 01, 2026

Regular, Invited, Tutorial
& Open Invited Track Nov. 02, 2026

- Submission site already open
- More info: <https://conferences.ifac-control.org/sysid2027/>
- ... or just ask us!

Recent topics in system identification: learning models for decision and control

- Special Issue from IFAC SYSID 2027

Submission deadline: **01 October 2026**

This special issue highlights leading papers from the IFAC SYSID 2027 Symposium

Executive guest editor: **Dario Piga**



IFAC Journal of Systems and Control