

Meta Learning for System Identification

Marco Forgione

Dalle Molle Institute for Artificial Intelligence, USI-SUPSI, Lugano, Switzerland
Leuven, June 19, 2026



SUPSI

Motivation


- Machine learning models can effectively describe a wide variety of dynamical systems
- However, training is computationally expensive and data-hungry
- Can we **automate the design of algorithms**, model structures, optimizers **for a class** of system identification problems?
- This is meta learning, aka learning to learn!

The “sines” meta learning toy example

- Regression datasets $x \rightarrow y$: sines with randomized phase and amplitude

$$y(x) = A \sin(x + \psi)$$

- Simple structure, but just $K = 5$ points per dataset
- Would be trivial with $K = 50$ points per dataset, or if we knew the true model structure



Example from:


Finn, Chelsea, Pieter Abbeel, and Sergey Levine. "Model-agnostic meta-learning for fast adaptation of deep networks." International conference on machine learning. PMLR, 2017.

The “sines” meta learning toy example

- With a black-box structure (e.g., a small neural net) we overfit the $K = 5$ data points
- If we knew the structure, $K = 2$ data points would be enough to find A, ψ

$$y(x) = A \sin(x + \psi)$$


- What if we observe $K = 5$ points from many datasets? Can we learn the pattern?
- Results with existing meta learning approaches



From standard to meta system identification

- System Identification
 - **Dataset:** I/O trajectory $D = (u_{1:N}, y_{1:N})$ of one **particular** system
 - **Objective:** estimate a model from D

$$\theta = \text{Alg}(D)$$
$$\hat{y}_{1:N}^* = M_{\theta}(u_{1:N}^*)$$



- Alg is fixed (or *manually* tuned)

- Meta System Identification
 - **Meta dataset:** collection of datasets $\mathfrak{D} = \{D^{(1)}, D^{(2)}, \dots\}$ from **similar** systems
 - **Objective:** learn to make prediction on all systems in $\mathfrak{D}$

model-based meta learning


$$\hat{y}_{1:N}^* = M_{\text{Alg}_{\varphi}(D)}(u_{1:N}^*)$$

- Alg **is learned!**




in-context learning

$$\hat{y}_{1:N}^* = \mathcal{M}_{\varphi}(u_{1:N}^*, D)$$



In-context system identification

- Particular dynamical system:
 - Input/output **context** data, input **query** → predicted query output:



Forgione, M., Pura, F., & Piga, D. (2023). From system models to class models: An in-context learning paradigm. *IEEE Control Systems Letters*, 7, 3513-3518.

Large Language Model analogy

Output (Claude Sonnet 4.6)

- **Grote Markt & Town Hall** — the obvious starting point. The 15th-century Gothic Town Hall is one of the most extravagantly decorated secular buildings in all of Belgium: every surface bristles with carved statues and tracery.
- **Saint Peter's Church** — don't just glance at the exterior. Inside you'll find major works by Dierick Bouts, the 15th-century Flemish master who spent most of his career in Leuven. His *Last Supper* altarpiece is here. There's also a small crypt worth descending into.






Context

A user is passionate about history and is planning a trip to Leuven.


Query

Can you recommend some places to visit?



Meta model architecture

- Meta-model extended to process the query's input/output initial condition
- Parameterized as a Transformer



Meta model training

STEPS

1. Collect a meta dataset of similar systems from a **simulator**:

$$\mathcal{D} = (D^{(1)}, D^{(2)}, \dots), \quad i=1, 2, \dots, \infty$$

2. Meta training **offline** to estimate a meta-model $\mathcal{M}_\phi$:

$$\hat{\phi} = \arg \min_{\phi} \sum_i \text{MSE}(D^{(i)}, \mathcal{M}_\phi)$$

3. Make predictions on new (real) system, in a **zero-shot** fashion (no training):

$$\hat{\mathbf{y}}^* = \mathcal{M}_{\hat{\phi}}(\mathbf{u}^*, \mathbf{D})$$

- Technically, it is just a supervised learning problem – standard training.
- The trained $\mathcal{M}$ may be seen as a **learned** identification + simulation algorithm...
- ... or as a *specialized* question answering system – a kind of SYSID GPT.

Meta Training – System class

Wiener-Hammerstein Systems





- Sequential LTI $\rightarrow F(\cdot) \rightarrow$ LTI
- Random LTI, order ≤ 5
- $F(\cdot)$: random feed-forward NN

- Meta training on **~100 M simulated** Wiener-Hammerstein systems
- The input $u_{1:N}$ is a multisine signal with random phase and spectrum
- White noise added to the simulated output
- Transformer model with **~5 M parameters**, training time **~1 day** on an Nvidia 4090 GPU...

Testing – Nonlinear benchmarks (www.nonlinearbenchmark.org)

- Real Wiener-Hammerstein benchmark





RMSE * 1000	Inference time (s)
1.4	5.91

State-of-the-art

Rufolo, M., Piga, D., Pura, F., & Forgone, M. (2025). In-Context Learning for System Identification: Recent Advances. Submitted to a journal, available [here](#).

Testing – Nonlinear benchmarks

- Coupled Electric Drives (CED)





RMSE	Inference time (s)
[0.068,0.055]	0.17

State-of-the-art

Rufolo, M., Piga, D., Pura, F., & Forgione, M. (2025). In-Context Learning for System Identification: Recent Advances. Submitted to a journal, available [here](#).

Testing – Nonlinear benchmarks

- Cascaded tanks with Overflow






RMSE	Inference time (s)
0.429	0.22

State-of-the-art

Rufolo, M., Piga, D., Pura, F., & Forgiione, M. (2025). In-Context Learning for System Identification: Recent Advances. Submitted to a journal, available [here](#).

Other objectives: meta state estimation

- A meta state estimator that generalizes to different systems from the context



- Experiments on speed estimation of BLDC motors

Colombo, A., Busetto, R., Breschi, V., Forgione, M., Piga, D., & Formentin, S. (2025). In-Context Learning for Zero-Shot Speed Estimation of BLDC motors. To be presented at IFAC World Congress 2026.


Busetto, R., Breschi, V., Forgione, M., Piga, D., & Formentin, S. (2024). In-context learning of state estimators. Presented at SYSID 2024.

Whitening the meta-learning box

The meta-learning grayscale

model-based meta learning

$$\hat{y}_{1:N}^* = M_{\text{Alg}_{\varphi}(D)}(u_{1:N}^*)$$




in-context learning

$$\hat{y}_{1:N}^* = \mathcal{M}_{\varphi}(u_{1:N}^*, D)$$

- In-context learning: prediction is a **black-box function** of dataset and new data.
- Other extreme: **learn hyper-parameters** φ of algorithm Alg_{φ} that returns (interpretable) **model parameters**
 - Learn the learning rate, initial value for gradient-based optimization (MAML), etc

Meta state estimation: whitening the box

- Transformer returns estimator gain K + **physical** parameters θ instead of predictions directly.



Model-based meta learning



In-context learning

- On-going work with Université de Poitiers (Hugo Koide, Guillaume Mercère)

Conclusions

Takeaways:

- In-context system identification, estimation, & control
- Off-line: learn how to map data into prediction
- On-line: provide context data, query input → query output ([zero shot](#))

Current activities:

- Real-world applications
- Whitening the meta-learning box
- Providing guarantees
- Algorithm discovery with LLM-driven agents