

Meta Learning for System Identification

Marco Forgione

Dalle Molle Institute for Artificial Intelligence, USI-SUPSI, Lugano, Switzerland
Leuven, June 19, 2026



SUPSI

Motivation


- Machine learning models can effectively describe a wide variety of dynamical systems
- However, training is computationally expensive and data-hungry
- Can we **automate the design of algorithms**, model structures, optimizers **for a class** of system identification problems?
- This is meta learning, aka learning to learn!

The “sines” meta learning toy example

- Regression datasets $x \rightarrow y$: sines with randomized phase and amplitude

$$y(x) = A \sin(x + \psi)$$

- Simple structure, but just $K = 5$ points per dataset
- Would be trivial with $K = 50$ points per dataset, or if we knew the true model structure



Example from:


Finn, Chelsea, Pieter Abbeel, and Sergey Levine. "Model-agnostic meta-learning for fast adaptation of deep networks." International conference on machine learning. PMLR, 2017.

The “sines” meta learning toy example

- With a black-box structure (e.g., a small neural net) we overfit the $K = 5$ data points
- If we knew the structure, $K = 2$ data points would be enough to find A, ψ

$$y(x) = A \sin(x + \psi)$$


- What if we observe $K = 5$ points from many datasets? Can we learn the pattern?
- Results with existing meta learning approaches



From standard to meta system identification

- System Identification
 - **Dataset:** I/O trajectory $D = (u_{1:N}, y_{1:N})$ of one **particular** system
 - **Objective:** estimate a model from D

$$\theta = \text{Alg}(D)$$
$$\hat{y}_{1:N}^* = M_{\theta}(u_{1:N}^*)$$



- Alg is fixed (or *manually* tuned)

- Meta System Identification
 - **Meta dataset:** collection of datasets $\mathfrak{D} = \{D^{(1)}, D^{(2)}, \dots\}$ from **similar** systems
 - **Objective:** learn to make prediction on all systems in $\mathfrak{D}$

model-based meta learning


$$\hat{y}_{1:N}^* = M_{\text{Alg}_{\varphi}(D)}(u_{1:N}^*)$$

- Alg **is learned!**




in-context learning

$$\hat{y}_{1:N}^* = \mathcal{M}_{\varphi}(u_{1:N}^*, D)$$



In-context system identification

- Particular dynamical system:
 - Input/output **context** data, input **query** → predicted query output:



Forgione, M., Pura, F., & Piga, D. (2023). From system models to class models: An in-context learning paradigm. *IEEE Control Systems Letters*, 7, 3513-3518.

Large Language Model analogy

Output (Claude Sonnet 4.6)

- **Grote Markt & Town Hall** — the obvious starting point. The 15th-century Gothic Town Hall is one of the most extravagantly decorated secular buildings in all of Belgium: every surface bristles with carved statues and tracery.
- **Saint Peter's Church** — don't just glance at the exterior. Inside you'll find major works by Dierick Bouts, the 15th-century Flemish master who spent most of his career in Leuven. His *Last Supper* altarpiece is here. There's also a small crypt worth descending into.






Context

A user is passionate about history and is planning a trip to Leuven.


Query

Can you recommend some places to visit?



Meta model architecture

- Meta-model extended to process the query's input/output initial condition
- Parameterized as a Transformer



Meta model training

STEPS

1. Collect a meta dataset of similar systems from a **simulator**:

$$\mathcal{D} = (D^{(1)}, D^{(2)}, \dots), \quad i=1, 2, \dots, \infty$$

2. Meta training **offline** to estimate a meta-model $\mathcal{M}_\phi$:

$$\hat{\phi} = \arg \min_{\phi} \sum_i \text{MSE}(D^{(i)}, \mathcal{M}_\phi)$$

3. Make predictions on new (real) system, in a **zero-shot** fashion (no training):

$$\hat{\mathbf{y}}^* = \mathcal{M}_{\hat{\phi}}(\mathbf{u}^*, \mathbf{D})$$

- Technically, it is just a supervised learning problem – standard training.
- The trained $\mathcal{M}$ may be seen as a **learned** identification + simulation algorithm...
- ... or as a *specialized* question answering system – a kind of SYSID GPT.

Meta Training – System class

Wiener-Hammerstein Systems





- Sequential LTI $\rightarrow F(\cdot) \rightarrow$ LTI
- Random LTI, order ≤ 5
- $F(\cdot)$: random feed-forward NN

- Meta training on **~100 M simulated** Wiener-Hammerstein systems
- The input $u_{1:N}$ is a multisine signal with random phase and spectrum
- White noise added to the simulated output
- Transformer model with **~5 M parameters**, training time **~1 day** on an Nvidia 4090 GPU...

Testing – Nonlinear benchmarks (www.nonlinearbenchmark.org)

- Real Wiener-Hammerstein benchmark





RMSE * 1000	Inference time (s)
1.4	5.91

State-of-the-art

Rufolo, M., Piga, D., Pura, F., & Forgone, M. (2025). In-Context Learning for System Identification: Recent Advances. Submitted to a journal, available [here](#).

Testing – Nonlinear benchmarks

- Coupled Electric Drives (CED)





RMSE	Inference time (s)
[0.068,0.055]	0.17

State-of-the-art

Rufolo, M., Piga, D., Pura, F., & Forgione, M. (2025). In-Context Learning for System Identification: Recent Advances. Submitted to a journal, available [here](#).

Testing – Nonlinear benchmarks

- Cascaded tanks with Overflow






RMSE	Inference time (s)
0.429	0.22

State-of-the-art

Rufolo, M., Piga, D., Pura, F., & Forgone, M. (2025). In-Context Learning for System Identification: Recent Advances. Submitted to a journal, available [here](#).

Other objectives: meta state estimation

- A meta state estimator that generalizes to different systems from the context



- Experiments on speed estimation of BLDC motors

Colombo, A., Busetto, R., Breschi, V., Forgone, M., Piga, D., & Formentin, S. (2025). In-Context Learning for Zero-Shot Speed Estimation of BLDC motors. To be presented at IFAC World Congress 2026.


Busetto, R., Breschi, V., Forgone, M., Piga, D., & Formentin, S. (2024). In-context learning of state estimators. Presented at SYSID 2024.

Whitening the meta-learning box

The meta-learning grayscale

model-based meta learning

$$\hat{y}_{1:N}^* = M_{\text{Alg}_{\varphi}(D)}(u_{1:N}^*)$$




in-context learning

$$\hat{y}_{1:N}^* = \mathcal{M}_{\varphi}(u_{1:N}^*, D)$$

- In-context learning: prediction is a **black-box function** of dataset and new data.
- Other extreme: **learn hyper-parameters** φ of algorithm Alg_{φ} that returns (interpretable) **model parameters**
 - Learn the learning rate, initial value for gradient-based optimization (MAML), etc

Meta state estimation: whitening the box

- Transformer returns estimator gain K + **physical** parameters θ instead of predictions directly.



Model-based meta learning



In-context learning

- On-going work with Université de Poitiers (Hugo Koide, Guillaume Mercère)

Conclusions

Takeaways:

- In-context system identification, estimation, & control
- Off-line: learn how to map data into prediction
- On-line: provide context data, query input → query output ([zero shot](#))

Current activities:

- Real-world applications
- Whitening the meta-learning box
- Providing guarantees
- Algorithm discovery with LLM-driven agents



INTERNATIONAL FEDERATION
OF AUTOMATIC CONTROL



21st IFAC Symposium on System Identification SysID2027

July 7-9 2027, Lyon, France



International Program Committee

IPC Chairs

Maarten Schoukens,
Marco Forgione

Editor

Kévin Colin

National Organizing Committee

NOC Chairs

Xavier Bombois
Marion Gilson,
Guillaume Mercère

Important Dates (Deadlines)

Submission site opens Jun. 01, 2026

Joint journal/conference
submission papers Oct. 01, 2026

Invited session,
Open Invited Track
& Tutorial proposals Oct. 01, 2026

Regular, Invited, Tutorial
& Open Invited Track Nov. 02, 2026

- Submission site already open
- More info: <https://conferences.ifac-control.org/sysid2027/>
- ... or just ask us!

**Recent topics in system identification: learning models for
decision and control**

- Special Issue from IFAC SYSID 2027

Submission deadline: **01 October 2026**

This special issue highlights leading papers from the IFAC SYSID
2027 Symposium

Executive guest editor: **Dario Piga**



IFAC Journal of Systems and Control